

# JULIE PHAM - 121540249 - Final Exam

## ----- Question 1 -----

In [209...]

```
# 1.1 Import the mtcars dataset and check data types
import pandas as pd
import numpy as np

# Load dataset
df = pd.read_csv("mtcars.csv")
df.head()
```

Out[209...]

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

In [210...]

```
print("Data Types")
df.dtypes
```

Data Types

Out[210...]

```
model      object
mpg       float64
cyl        int64
disp      float64
hp         int64
drat      float64
wt         float64
qsec      float64
vs          int64
am          int64
gear      int64
carb      int64
dtype: object
```

In [211...]

```
# 1.2. Are there any missing values in the dataset? --> NO
print("Missing Values")
df.isnull().sum()
```

Missing Values

Out[211...]

```
model      0
mpg       0
cyl       0
disp      0
hp        0
drat      0
wt        0
qsec      0
vs        0
am        0
gear      0
carb      0
dtype: int64
```

In [212...]

```
# 1.3. Print the descriptive statistics of the mtcars data to understand the data a little better (min, max, mean, median, 1st and
print(df.describe().T)
```

	count	mean	std	min	25%	50%	75%	\
mpg	32.0	20.090625	6.026948	10.400	15.42500	19.200	22.80	
cyl	32.0	6.187500	1.785922	4.000	4.00000	6.000	8.00	
disp	32.0	230.721875	123.938694	71.100	120.82500	196.300	326.00	
hp	32.0	146.687500	68.562868	52.000	96.50000	123.000	180.00	
drat	32.0	3.596563	0.534679	2.760	3.08000	3.695	3.92	
wt	32.0	3.217250	0.978457	1.513	2.58125	3.325	3.61	
qsec	32.0	17.848750	1.786943	14.500	16.89250	17.710	18.90	
vs	32.0	0.437500	0.504016	0.000	0.00000	0.000	1.00	
am	32.0	0.406250	0.498991	0.000	0.00000	0.000	1.00	
gear	32.0	3.687500	0.737804	3.000	3.00000	4.000	4.00	
carb	32.0	2.812500	1.615200	1.000	2.00000	2.000	4.00	
								max
mpg		33.900						
cyl		8.000						
disp		472.000						
hp		335.000						
drat		4.930						
wt		5.424						
qsec		22.900						
vs		1.000						
am		1.000						
gear		5.000						
carb		8.000						

In [213...]

```
# 1.4. Splitting the Data-Set into Independent and Dependent Features
X = df.drop(columns=["am"]) # Independent variables
y = df["am"] # Dependent variable Transmission(am))
```

In [214...]

```
# Count the occurrences ---> not imbalance
counts = y.value_counts()
# Calculate percentages
percentages = counts / len(y) * 100
print("Counts:\n", counts)
print("\nPercentages:\n", percentages)
```

Counts:

```
am
0    19
1    13
Name: count, dtype: int64
```

Percentages:

```
am
0    59.375
1    40.625
Name: count, dtype: float64
```

In [215...]

```
# 1.5. Convert categorical variable into numeric Using one hot encoding method
from sklearn.preprocessing import OneHotEncoder, StandardScaler
categorical_cols = X.select_dtypes(include=['object']).columns
if len(categorical_cols) > 0:
    encoder = OneHotEncoder(drop='first', sparse_output=False)
    encoded_df = pd.DataFrame(
        encoder.fit_transform(X[categorical_cols]),
        columns=encoder.get_feature_names_out(categorical_cols),
        index=X.index
    )
    X = pd.concat([X.drop(columns=categorical_cols), encoded_df], axis=1)
X.head()
```

Out[215...]

	mpg	cyl	disp	hp	drat	wt	qsec	vs	gear	carb	...	model_Merc_280C	model_Merc_450SE	model_Merc_450SL	model_Merc_450SLC	model_Pontiac_Firebird	model_F
0	21.0	6	160.0	110	3.90	2.620	16.46	0	4	4	...	0.0	0.0	0.0	0.0	0.0	
1	21.0	6	160.0	110	3.90	2.875	17.02	0	4	4	...	0.0	0.0	0.0	0.0	0.0	
2	22.8	4	108.0	93	3.85	2.320	18.61	1	4	1	...	0.0	0.0	0.0	0.0	0.0	
3	21.4	6	258.0	110	3.08	3.215	19.44	1	3	1	...	0.0	0.0	0.0	0.0	0.0	
4	18.7	8	360.0	175	3.15	3.440	17.02	0	3	2	...	0.0	0.0	0.0	0.0	0.0	

5 rows × 41 columns



In [216...]

```
# 1.6. Normalize the data set.
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

In [217...]

```
# 1.7. Divide the dataset to training and test sets.
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)
```

In [218...]

```
# 1.8. Use the decision tree, KNN to predict the test set out values.
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
```

In [219...]

```
# Decision Tree
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
```

In [220...]

```
# KNN
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
```

In [221...]

```
# 1.9. Display the confusion matrix to evaluate the model performance.
from sklearn.metrics import confusion_matrix, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt
```

```

cm_dt = confusion_matrix(y_test, y_pred_dt)
cm_knn = confusion_matrix(y_test, y_pred_knn)


plt.figure(figsize=(12, 5))

# Decision Tree heatmap
plt.subplot(1, 2, 1)
sns.heatmap(cm_dt, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title("Decision Tree - Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")

# KNN heatmap
plt.subplot(1, 2, 2)
sns.heatmap(cm_knn, annot=True, fmt="d", cmap="Greens", cbar=False)
plt.title("KNN - Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")

plt.tight_layout()
plt.show()

```



In [222...]

```

print("Confusion Matrix - Decision Tree")
print(cm_dt)

```

```
print("Confusion Matrix - KNN")
print(cm_knn)
```

```
Confusion Matrix - Decision Tree
[[4 0]
 [1 2]]
Confusion Matrix - KNN
[[4 0]
 [1 2]]
```

Both Decision Tree and KNN produced the exact same confusion matrix, meaning they made the same predictions on this dataset. The models are very good at predicting class 0 (automatic), but slightly weaker at predicting class 1 (manual).

In [224...]

```
#1.10. Evaluate the model performance by computing Accuracy.
print("Accuracy Scores")
print(f"Decision Tree Accuracy: {accuracy_score(y_test, y_pred_dt):.2f}")
print(f"KNN Accuracy: {accuracy_score(y_test, y_pred_knn):.2f}")
```

```
Accuracy Scores
Decision Tree Accuracy: 0.86
KNN Accuracy: 0.86
```

In this case, both the Decision Tree and KNN models yielded the exact same performance metrics on this dataset. The result show that they have similar predictive power in this specific scenario (classify transmission is automatic or manual).

## ----- Question 2 -----

In [227...]

```
# 2.1. Import the Iris dataset
import pandas as pd

# Load dataset
df1 = pd.read_csv("iris_csv.csv")
df1.head()
```

Out[227...]

	sepallength	sepalwidth	petallength	petalwidth	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [228...]

```
#2.2. Remove the Species column from the iris dataset?
X = df1.drop(columns=["class"]) # Keep features only
y_true = df1["class"] # Keep the actual species for comparison
X.head()
```

Out[228...]

	sepallength	sepalwidth	petallength	petalwidth
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

In [229...]

```
#2.3. Create a model with three clusters that can cluster the observations in the data
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
kmeans.fit(X)
clusters = kmeans.labels_
```

C:\Users\ADMIN\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.  
 warnings.warn(

In [230...]

```
# 2.4 Compare clusters with actual species
comparison_df = pd.DataFrame({"Actual": y_true, "Cluster": clusters})
print("First 10 cluster assignments")
print(comparison_df.head(10))
```

```
First 10 cluster assignments
    Actual Cluster
0  Iris-setosa      1
1  Iris-setosa      1
2  Iris-setosa      1
3  Iris-setosa      1
4  Iris-setosa      1
5  Iris-setosa      1
6  Iris-setosa      1
7  Iris-setosa      1
8  Iris-setosa      1
9  Iris-setosa      1
```

In [231...]

```
# Map cluster numbers to species using value_counts
mapping = {}
for cluster in range(3):
    mask = clusters == cluster
    mapping[cluster] = y_true[mask].value_counts().idxmax()
# finds the species with the highest count in that cluster.
# Predicted species based on mapping
predicted_species = [mapping[cluster] for cluster in clusters]

# Count mistakes
mistakes = np.sum(predicted_species != y_true)
print(f"Number of mistakes made by the model: {mistakes} out of {len(y_true)}")
```

Number of mistakes made by the model: 16 out of 150

In [232...]

```
16/150*100
```

Out[232...]

```
10.666666666666668
```

**16 mistakes out of 150 samples, which is roughly 10.7% error.** That is pretty good for an unsupervised method like KMeans.

In [234...]

```
# Create comparison DataFrame
comparison_df = pd.DataFrame({
    "Actual Species": y_true,
    "Predicted Species": predicted_species
})
# Filter only incorrect predictions
incorrect_df = comparison_df[comparison_df["Actual Species"] != comparison_df["Predicted Species"]]
print("Incorrect Predictions")
incorrect_df
```

Incorrect Predictions

Out[234...]

	Actual Species	Predicted Species
52	Iris-versicolor	Iris-virginica
77	Iris-versicolor	Iris-virginica
101	Iris-virginica	Iris-versicolor
106	Iris-virginica	Iris-versicolor
113	Iris-virginica	Iris-versicolor
114	Iris-virginica	Iris-versicolor
119	Iris-virginica	Iris-versicolor
121	Iris-virginica	Iris-versicolor
123	Iris-virginica	Iris-versicolor
126	Iris-virginica	Iris-versicolor
127	Iris-virginica	Iris-versicolor
133	Iris-virginica	Iris-versicolor
138	Iris-virginica	Iris-versicolor
142	Iris-virginica	Iris-versicolor
146	Iris-virginica	Iris-versicolor
149	Iris-virginica	Iris-versicolor

### Most errors are between Versicolor and Virginica

Cluster assignments for these two species overlap a lot.

KMeans groups by feature similarity, and in the iris dataset, Versicolor and Virginica features are closer to each other compared to Setosa.

Setosa is never misclassified, which is expected because Setosa is well-separated in feature space (it's easily clustered).

In [265...]

```
from sklearn.metrics import classification_report

# Generate classification report
report = classification_report(y_true, predicted_species)
print("Classification Report")
print(report)
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	50
Iris-versicolor	0.77	0.96	0.86	50
Iris-virginica	0.95	0.72	0.82	50
accuracy			0.89	150
macro avg	0.91	0.89	0.89	150
weighted avg	0.91	0.89	0.89	150

**Accuracy = 0.89 (89%)** --> Good

Setosa: Perfectly separated; Setosa is easily clustered.

Versicolor: Most Versicolor points were correctly identified; a few were misclassified as Virginica (lower precision).

Virginica: Many Virginica points were misclassified as Versicolor (lower recall).



In [236...]

```
# Visualization: Clusters vs Actual Species
plt.figure(figsize=(12, 5))

# Left: KMeans clusters
plt.subplot(1, 2, 1)
sns.scatterplot(x=X.iloc[:, 0], y=X.iloc[:, 1], hue=clusters, palette="Set1", s=60)
plt.title("KMeans Clusters")
plt.xlabel(X.columns[0])
plt.ylabel(X.columns[1])
plt.legend(title="Cluster", loc="best")

# Right: True species
plt.subplot(1, 2, 2)
sns.scatterplot(x=X.iloc[:, 0], y=X.iloc[:, 1], hue=y_true, palette="Set2", s=60)
plt.title("Actual Species")
plt.xlabel(X.columns[0])
plt.ylabel(X.columns[1])
plt.legend(title="Species", loc="best")

plt.tight_layout()
plt.show()
```



Setosa points are isolated, while Versicolor and Virginica points are partially overlapping, explaining these 16 errors.