

```
In [6]: !pip install pyreadstat  
import pyreadstat as pyd
```

```
Requirement already satisfied: pyreadstat in c:\users\ryan\anaconda3\lib\site-packages (1.2.8)  
Requirement already satisfied: pandas>=1.2.0 in c:\users\ryan\anaconda3\lib\site-packages (from pyreadstat) (2.2.2)  
Requirement already satisfied: numpy>=1.26.0 in c:\users\ryan\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (1.2  
6.4)  
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\ryan\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat)  
(2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in c:\users\ryan\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (202  
4.1)  
Requirement already satisfied: tzdata>=2022.7 in c:\users\ryan\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (20  
23.3)  
Requirement already satisfied: six>=1.5 in c:\users\ryan\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas>=1.2.  
0->pyreadstat) (1.16.0)
```

```
In [10]: df1, meta = pyd.read_sas7bdat('./retention.sas7bdat')
```

Q1

```
In [13]: # Display the first few rows  
print(df1.head())
```

	Att_hrs_fall	GENDER	HIGH SCHOOL_PERCENTILE	AGE	Att_hrs_spr	\
0	13.0	F		97.0	18.910335	15.0
1	14.0	M		97.0	18.036961	15.0
2	15.0	M		83.0	18.162902	15.0
3	14.0	M		92.0	18.590007	12.0
4	16.0	F		80.0	18.880219	16.0

	Avg_income	Distance	Dropped_course	Major_rate	SAT	...	Legacynum	\
0	30573.0	5.083094		0	0.821782	1160.0	...	0
1	27305.0	2.198764		0	0.795455	1050.0	...	0
2	30573.0	5.083094		0	0.843750	1140.0	...	1
3	35865.0	3.245574		0	NaN	1090.0	...	0
4	40125.0	19.486585		0	0.840391	1150.0	...	0

	Stu_worker_ind	Need_pct_met	Perc_hrs_comp_fall	Hs_rate	Dorm_rate	\
0	0	1.000000		1.0	0.812500	0.8
1	0	0.977366		1.0	1.000000	0.8
2	0	0.913965		1.0	0.812500	0.8
3	1	1.000000		1.0	0.847826	0.8
4	1	1.000000		1.0	0.875000	0.8

	Instate	Transcrip	Fall_GPA	Target
0	1	0	3.307692	0
1	1	1	2.821429	0
2	1	0	3.133333	0
3	1	0	2.678571	0
4	1	0	2.500000	0

[5 rows x 21 columns]

In [15]:

```
# Display metadata
print("Column Names:", meta.column_names)
print("Column Labels:", meta.column_labels)
```

Column Names: ['Att_hrs_fall', 'GENDER', 'HIGH SCHOOL_PERCENTILE', 'AGE', 'Att_hrs_spr', 'Avg_income', 'Distance', 'Dropped_course', 'Major_rate', 'SAT', 'Extra_curr', 'Legacynum', 'Stu_worker_ind', 'Need_pct_met', 'Perc_hrs_comp_fall', 'Hs_rate', 'Dorm_rate', 'Instate', 'Transcrip', 'Fall_GPA', 'Target']
Column Labels: ['Att_hrs_fall', 'GENDER', 'HIGH SCHOOL_PERCENTILE', 'AGE', 'Att_hrs_spr', 'Avg_income', 'Distance', 'Dropped_course', 'Major_rate', 'SAT', 'Extra_curr', 'Legacynum', 'Stu_worker_ind', 'Need_pct_met', 'Perc_hrs_comp_fall', 'Hs_rate', 'Dorm_rate', 'Instate', 'Transcrip', 'Fall_GPA', 'Target']

```
In [17]: #Q1a - How many observations (rows) exist in the Retention table?
```

```
print(f"Number of observations (rows): {df1.shape[0]}")
```

```
Number of observations (rows): 2626
```

```
In [19]: #Q1b - • How many variables (columns) exist in the Retention table?
```

```
print(f"Number of variables (columns): {df1.shape[1]}")
```

```
Number of variables (columns): 21
```

Q1c - How are the initial roles and levels of the variables determined? Can we change the default settings? Explain by example.

Explanation of roles and levels

Variables roles are categorized as input or target based on their function in the analysis. Input variables are the independent and used to determine the target while target variables are dependent and is the result

Variable levels are nominal, ordinal and interval/ratio

We can override their default setting if required

```
In [22]: # Example: Convert AGE into categorical age groups
```

```
import pandas as pd
# check if its numerical or categorical
print(df1['AGE'].dtype)
```

```
float64
```

```
In [24]: # change it from numerical to categorical
```

```
df1['AGE'] = df1['AGE'].astype('category')
print(df1['AGE'].dtype)
```

category

Q2

```
In [27]: #check for missing values  
print(df1.isnull().sum())
```

```
Att_hrs_fall          0  
GENDER                0  
HIGH SCHOOL PERCENTILE 263  
AGE                   0  
Att_hrs_spr           0  
Avg_income            157  
Distance              98  
Dropped_course        0  
Major_rate             79  
SAT                   0  
Extra_curr             0  
LegacyCnum             0  
Stu_worker_ind         0  
Need_pct_met           0  
Perc_hrs_comp_fall    0  
Hs_rate                805  
Dorm_rate              0  
Instate                0  
Transcrip              0  
Fall_GPA                0  
Target                  0  
dtype: int64
```

No, variables with missing values are not rejected by Python. Avg_income has 157 missing variables, HIGH SCHOOL PERCENTILE has 263 missing variables and Hs_rate has 805 missing variables.

```
In [30]: # check which numeric variables are assigned as nominal variables  
variables = df1.select_dtypes(include=['object'])  
print(variables.columns)
```

```
Index(['GENDER', 'Dropped_course', 'Legacynum', 'Stu_worker_ind', 'Instate',
       'Transcrip', 'Target'],
      dtype='object')
```

Usually variables with the data type 'object' are typically treated as nominal variables, sometimes numeric-looking values can also be stored as objects sucha as (0,1) this can be the case if their data was stored as string format originally.

Is there a maximum number of levels for character variables

There is no maximum number of levels for character variables a column can have multiple unique variables, however using too many unique variables can lead to technical challenges and can make the variable less useful during analysis.

Q3

```
In [43]: #How many input variables in the DataFrame? How many target variables? How many of the input variables are nominal? Are there
#Total input variables: 20
#Target variables: 1 (Target)
#Nominal (categorical) input variables: 2
#'Att_hrs_fall'
#'GENDER'
#Rejected variables: None (no columns were discarded or found irrelevant based on this analysis)
```

```
In [ ]: #What is the sample size?
# Total records: 2626. 10% sample (commonly used in exploratory data analysis): 262 rows
# This is a typical and manageable size for quick testing, modeling, or visualization.
```


```
In [49]: import matplotlib.pyplot as plt
```

```
In [51]: # Step 2: Plot bar chart of GENDER
gender_counts = df1['GENDER'].value_counts()
gender_percentages = gender_counts / gender_counts.sum() * 100

# Bar plot
```

```
plt.figure(figsize=(6, 4))
gender_counts.plot(kind='bar', color=['skyblue', 'lightcoral'])
plt.title('Gender Distribution')
plt.xlabel('Gender')
plt.ylabel('Number of Students')
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.tight_layout()
plt.show()

# Step 3: Calculate percentage of female students
female_percentage = gender_percentages.get('F', 0)
print(f"Percentage of female students: {female_percentage:.2f}%")
```



Percentage of female students: 60.24%

In [59]: `import pandas as pd
import matplotlib.pyplot as plt`

```
import numpy as np

# Step 1: Load and parse the CSV
file_path = "retention.csv"
with open(file_path, 'r') as f:
    lines = f.readlines()


header = lines[0].strip().split(",")
data = [line.strip().split(",") for line in lines[1:]]
df = pd.DataFrame(data, columns=[col.strip('') for col in header])

# Step 2: Count total and target=1 by gender
total_counts = df['GENDER'].value_counts()
target1_counts = df[df['Target'] == '1']['GENDER'].value_counts()

# Step 3: Align bars side-by-side
labels = sorted(total_counts.index)
x = np.arange(len(labels)) # Label locations
width = 0.35 # width of the bars

# Step 4: Plot
fig, ax = plt.subplots(figsize=(6, 4))
bars1 = ax.bar(x - width/2, [total_counts.get(label, 0) for label in labels], width, label='Total Students', color='lightgray')
bars2 = ax.bar(x + width/2, [target1_counts.get(label, 0) for label in labels], width, label='Target = 1', color='dodgerblue')

# Labels and formatting
ax.set_xlabel('Gender')
ax.set_ylabel('Number of Students')
ax.set_title('Gender Distribution with Highlight on Target=1')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()
ax.grid(axis='y')
plt.tight_layout()
plt.show()
```



```
In [61]: import pandas as pd
import matplotlib.pyplot as plt


# Step 1: Load and parse CSV again if not already done
file_path = "retention.csv"
with open(file_path, 'r') as f:
    lines = f.readlines()

header = lines[0].strip().split(",")
data = [line.strip().split(",") for line in lines[1:]]
df = pd.DataFrame(data, columns=[col.strip(' "') for col in header])

# Step 2: Convert Avg_income to numeric (handle non-numeric entries if any)
df['Avg_income'] = pd.to_numeric(df['Avg_income'], errors='coerce')

# Step 3: Plot histogram
```

```
plt.figure(figsize=(7, 5))
plt.hist(df['Avg_income'].dropna(), bins=20, color='mediumseagreen', edgecolor='black')
plt.title('Distribution of Average Income')
plt.xlabel('Avg_income')
plt.ylabel('Number of Students')
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```



In [65]: `import pandas as pd`

```
# Step 1: Load and parse CSV
file_path = "retention.csv"
with open(file_path, 'r') as f:
    lines = f.readlines()

header = lines[0].strip().split(",")
data = [line.strip().split(",") for line in lines[1:]]
df = pd.DataFrame(data, columns=[col.strip(' ') for col in header])

# Step 2: Get the last 10 rows
last_10_rows = df.tail(10)

# Step 3: Display the result
print(last_10_rows)
```

	Att_hrs_fall	GENDER	HIGH SCHOOL PERCENTILE	AGE	\
2616	"15.0	M	59.0	18.60643394934976	
2617	"12.0	M	13.0	18.88021902806297	
2618	"13.0	F	51.0	18.48870636550308	
2619	"14.0	M	62.0	18.28062970568104	
2620	"14.0	M	0.6	18.967830253251197	
2621	"13.0	M	42.0	18.151950718685832	
2622	"13.0	M	74.0	18.650239561943874	
2623	"16.0	F	0.6	18.102669404517453	
2624	"12.0	F	71.0	18.31895961670089	
2625	"16.0	F	95.0	18.6666666666666668	

	Att_hrs_spr	Avg_income	Distance	Dropped_course	\
2616	16.0	36760.0	156.74223825933137	0	
2617	18.0	26725.0	167.68172698652128	0	
2618	13.0	53190.0	1186.8518671692395	0	
2619	16.0	60898.0	1154.751207131066	0	
2620	13.0			0	
2621	12.0	48581.0	179.44467807371302	0	
2622	18.0		110.99359173168243	0	
2623	12.0	73510.0	356.6742856734745	0	
2624	13.0			0	
2625	16.0	59599.0	91.02294960168159	0	

	Major_rate	SAT	...	Legacynum	Stu_worker_ind	\
2616	0.8403908794788274	1290.0	...	0	1	
2617	0.7954545454545454	1060.0	...	0	0	
2618	0.8315847598012148	980.0	...	0	0	
2619	0.8315109343936382	1070.0	...	1	0	
2620	0.8292682926829268	1140.0	...	0	0	
2621	0.8383838383838383	1080.0	...	0	0	
2622	0.8315109343936382	1110.0	...	0	1	
2623	0.8285714285714286	1150.0	...	0	0	
2624	0.8292682926829268	1060.0	...	0	0	
2625	0.8295148247978437	1140.0	...	0	0	

	Need_pct_met	Perc_hrs_comp_fall	Hs_rate	Dorm_rate	Instate	\
2616	0.8166249293157767	0.06666666666666667	0.77381	0.817533	1	
2617	0.9697547419782209		0.5	0.6	0.845717	0
2618		1.0	0.7	0.6	0.82125	0
2619		0.0	0.21428571428571427	0.625	0.845717	0

2620	1.0	0.2857142857142857	0.7	0.7	1
2621	1.0	0.5	0.7	0.7	1
2622	0.9018721037998146	0.7692307692307693	0.666667	0.7	0
2623	1.0	0.6875	0.6	0.7	0
2624	1.0	0.5	0.5	0.7	0
2625	1.0	0.25	0.5	0.6	1

	Transcrip	Fall_GPA	Target
2616	1	0.071428571	1"
2617	1	0.7	1"
2618	1	0.7	1"
2619	1	0.692307692	1"
2620	0	0.285714286	1"
2621	1	1.346153846	1"
2622	1	1.0	1"
2623	1	0.7	1"
2624	1	1.0	1"
2625	0	0.78125	1"

[10 rows x 21 columns]

In []: