

```
In [93]: # Import necessary libraries
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [94]: #Load the dataset
```

```
spambase = pd.read_csv('../BAN 230/spambase_raw.csv')
spambase.head()
```

```
Out[94]:
```

	0	0.64	0.64.1	0.1	0.32	0.2	0.3	0.4	0.5	0.6	...	0.41	0.42	0.43	0.778	0.44	0.45	3.756	61	278	1
0	0.21	0.28	0.50	0.0	0.14	0.28	0.21	0.07	0.00	0.94	...	0.00	0.132	0.0	0.372	0.180	0.048	5.114	101	1028	1
1	0.06	0.00	0.71	0.0	1.23	0.19	0.19	0.12	0.64	0.25	...	0.01	0.143	0.0	0.276	0.184	0.010	9.821	485	2259	1
2	0.00	0.00	0.00	0.0	0.63	0.00	0.31	0.63	0.31	0.63	...	0.00	0.137	0.0	0.137	0.000	0.000	3.537	40	191	1
3	0.00	0.00	0.00	0.0	0.63	0.00	0.31	0.63	0.31	0.63	...	0.00	0.135	0.0	0.135	0.000	0.000	3.537	40	191	1
4	0.00	0.00	0.00	0.0	1.85	0.00	0.00	1.85	0.00	0.00	...	0.00	0.223	0.0	0.000	0.000	0.000	3.000	15	54	1

5 rows × 58 columns

```
In [95]: spambase.shape
```

```
Out[95]: (4600, 58)
```

```
In [96]: spambase.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 58 columns):
 #   Column  Non-Null Count  Dtype  
 --- 
  0   0        4600 non-null   float64
  1   0.64    4600 non-null   float64
  2   0.64.1  4600 non-null   float64
  3   0.1     4600 non-null   float64
  4   0.32    4600 non-null   float64
  5   0.2     4600 non-null   float64
  6   0.3     4600 non-null   float64
  7   0.4     4600 non-null   float64
  8   0.5     4600 non-null   float64
  9   0.6     4600 non-null   float64
  10  0.7     4600 non-null   float64
  11  0.64.2  4600 non-null   float64
  12  0.8     4600 non-null   float64
  13  0.9     4600 non-null   float64
  14  0.10    4600 non-null   float64
  15  0.32.1  4600 non-null   float64
  16  0.11    4600 non-null   float64
  17  1.29    4600 non-null   float64
  18  1.93    4600 non-null   float64
  19  0.12    4600 non-null   float64
  20  0.96    4600 non-null   float64
  21  0.13    4600 non-null   float64
  22  0.14    4600 non-null   float64
  23  0.15    4600 non-null   float64
  24  0.16    4600 non-null   float64
  25  0.17    4600 non-null   float64
  26  0.18    4600 non-null   float64
  27  0.19    4600 non-null   float64
  28  0.20    4600 non-null   float64
  29  0.21    4600 non-null   float64
  30  0.22    4600 non-null   float64
  31  0.23    4600 non-null   float64
  32  0.24    4600 non-null   float64
  33  0.25    4600 non-null   float64
  34  0.26    4600 non-null   float64
  35  0.27    4600 non-null   float64
  36  0.28    4600 non-null   float64
  37  0.29    4600 non-null   float64
  38  0.30    4600 non-null   float64
  39  0.31    4600 non-null   float64
  40  0.33    4600 non-null   float64
  41  0.34    4600 non-null   float64
  42  0.35    4600 non-null   float64
```

```
43  0.36    4600 non-null  float64
44  0.37    4600 non-null  float64
45  0.38    4600 non-null  float64
46  0.39    4600 non-null  float64
47  0.40    4600 non-null  float64
48  0.41    4600 non-null  float64
49  0.42    4600 non-null  float64
50  0.43    4600 non-null  float64
51  0.778   4600 non-null  float64
52  0.44    4600 non-null  float64
53  0.45    4600 non-null  float64
54  3.756   4600 non-null  float64
55  61      4600 non-null  int64
56  278     4600 non-null  int64
57  1       4600 non-null  int64
dtypes: float64(55), int64(3)
memory usage: 2.0 MB
```

```
In [97]: spambase.isnull().sum()
# there is no null in data
```

```
Out[97]: 0      0
          0.64   0
          0.64.1 0
          0.1    0
          0.32   0
          0.2    0
          0.3    0
          0.4    0
          0.5    0
          0.6    0
          0.7    0
          0.64.2 0
          0.8    0
          0.9    0
          0.10   0
          0.32.1 0
          0.11   0
          1.29   0
          1.93   0
          0.12   0
          0.96   0
          0.13   0
          0.14   0
          0.15   0
          0.16   0
          0.17   0
          0.18   0
          0.19   0
          0.20   0
          0.21   0
          0.22   0
          0.23   0
          0.24   0
          0.25   0
          0.26   0
          0.27   0
          0.28   0
          0.29   0
          0.30   0
          0.31   0
          0.33   0
          0.34   0
          0.35   0
          0.36   0
          0.37   0
          0.38   0
          0.39   0
```

```
0.40      0
0.41      0
0.42      0
0.43      0
0.778     0
0.44      0
0.45      0
3.756     0
61        0
278      0
1         0
dtype: int64
```

```
In [98]: #splitting feature X and target Y
X = spambase.iloc[:, :-1]
y = spambase.iloc[:, -1]
```

```
In [99]: # Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled
```

```
Out[99]: array([[ 3.45251829e-01,  5.19760953e-02,  4.35261246e-01, ...,
       -2.45283780e-03,  2.50545504e-01,  1.22818869e+00],
      [-1.45981828e-01, -1.64984012e-01,  8.51832832e-01, ...,
       1.45895187e-01,  2.22087495e+00,  3.25837649e+00],
      [-3.42475291e-01, -1.64984012e-01, -5.56575862e-01, ...,
       -5.21543111e-02, -6.24495382e-02, -1.52207080e-01],
      ...,
      [ 6.39992023e-01, -1.64984012e-01,  3.85264032e-02, ...,
       -1.19378942e-01, -2.36905791e-01, -2.72600020e-01],
      [ 2.80142011e+00, -1.64984012e-01, -5.56575862e-01, ...,
       -1.27478675e-01, -2.42036857e-01, -3.38568754e-01],
      [-3.42475291e-01, -1.64984012e-01,  7.32812379e-01, ...,
       -1.24232478e-01, -2.42036857e-01, -4.01239052e-01]])
```

```
In [100...]: # Split into train and test sets (70/30)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
```

```
In [101...]: # Train the SVM model
model = SVC(kernel='linear')
model.fit(X_train, y_train)
```

Out[101...]

SVC
SVC(kernel='linear')

In [102...]

```
#predict and model  
y_pred = model.predict(X_test)
```


In [103...]

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.91	0.94	0.93	803
1	0.92	0.88	0.90	577
accuracy			0.92	1380
macro avg	0.92	0.91	0.91	1380
weighted avg	0.92	0.92	0.91	1380

In [104...]

```
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix")  
plt.show()
```



Interpret

The SVM model with a linear kernel demonstrated strong performance in classifying emails as spam or not spam. It achieved an overall accuracy of 92%, as shown in the classification report.

From the confusion matrix, we observe: 758 legitimate (non-spam) emails were correctly identified. 45 legitimate emails were incorrectly classified as spam (false positives). 505 spam emails were correctly detected. 72 spam emails were missed and classified as non-spam (false negatives).

The model achieved a precision of 0.92 and recall of 0.88 for spam emails (label 1). This indicates it is highly accurate in labeling spam correctly but still misses some spam messages (the 72 false negatives). The F1-score of 0.90 for spam confirms a solid balance between "precision" and "recall". The false positives may arise from legitimate emails containing words or patterns resembling spam, while false negatives may occur when spam messages are too subtle or resemble normal communication. Overall, the model effectively minimizes both types of errors, making it a reliable spam filter with room for slight recall improvements on spam detection.

In [106...]

```
#80/20 train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```


In [107...]

```
# Train the SVM model
model = SVC(kernel='linear')
model.fit(X_train, y_train)
#predict and model
y_pred = model.predict(X_test)
#print the classification report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.92	0.95	0.93	530
1	0.93	0.88	0.91	390
accuracy			0.92	920
macro avg	0.92	0.92	0.92	920
weighted avg	0.92	0.92	0.92	920

In [108...]

```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



Compare the train-test split between 80/20 and 70/30:

In the 80/20 split scenario, the model trained on a larger portion of the data, with 80% used for training and 20% for testing. The model achieved the same overall accuracy of 92%, consistent with the 70/30 split. However, there were small but meaningful differences in the precision and recall values.

From the classification report: Spam (label 1) detection showed slightly higher precision (0.93 vs. 0.92) and same recall (0.88), leading to a small boost in F1-score (0.91 vs. 0.90). Non-spam (label 0) recall also improved (0.95 vs. 0.94), and precision increased to 0.92.

From the confusion matrix: False positives (non-spam predicted as spam) decreased from 45 to 24. False negatives (spam missed) slightly decreased from 72 to 47. The 80/20 split yielded slightly better classification performance, particularly in reducing misclassified legitimate emails (false positives). The model benefited from having more training data, improving generalization. The results suggest that, for this dataset size, using more data for training can lead to marginal gains in precision without sacrificing overall accuracy. This makes the 80/20 split a favorable option for deployment scenarios where spam precision is critical.

