

BAN 210: Workshop 1

Import data

```
In [1]: !pip install pyreadstat
```

```
Requirement already satisfied: pyreadstat in c:\users\admin\anaconda3\lib\site-packages (1.2.7)
Requirement already satisfied: pandas>=1.2.0 in c:\users\admin\anaconda3\lib\site-packages (from pyreadstat) (2.2.2)
Requirement already satisfied: numpy>=1.26.0 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (1.2
6.4)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadst
at) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (2024.
1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.2.0->pyreadstat) (202
3.3)
Requirement already satisfied: six>=1.5 in c:\users\admin\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas>=1.2.0
->pyreadstat) (1.16.0)
```

```
In [5]: import pyreadstat as pyds
df1, meta = pyds.read_sas7bdat('../BAN 210/retention.sas7bdat')
df1
```

Out[5]:

	Att_hrs_fall	GENDER	HIGH SCHOOL PERCENTILE	AGE	Att_hrs_spr	Avg_income	Distance	Dropped_course	Major_rate	SAT
0	13.0	F		97.0	18.910335	15.0	30573.0	5.083094	0	0.821782 1160.0
1	14.0	M		97.0	18.036961	15.0	27305.0	2.198764	0	0.795455 1050.0
2	15.0	M		83.0	18.162902	15.0	30573.0	5.083094	0	0.843750 1140.0
3	14.0	M		92.0	18.590007	12.0	35865.0	3.245574	0	NaN 1090.0
4	16.0	F		80.0	18.880219	16.0	40125.0	19.486585	0	0.840391 1150.0
...
2621	13.0	M		42.0	18.151951	12.0	48581.0	179.444678	0	0.838384 1080.0
2622	13.0	M		74.0	18.650240	18.0	NaN	110.993592	0	0.831511 1110.0
2623	16.0	F		0.6	18.102669	12.0	73510.0	356.674286	0	0.828571 1150.0
2624	12.0	F		71.0	18.318960	13.0	NaN	NaN	0	0.829268 1060.0
2625	16.0	F		95.0	18.666667	16.0	59599.0	91.022950	0	0.829515 1140.0

2626 rows × 21 columns

Q1:

How many observations (rows) exist in the Retention table? How many variables (columns) exist in the Retention table? How are the initial roles and levels of the variables determined? Can we change the default settings? Explain by example

In [13]: df1.shape

Out[13]: (2626, 21)

In [18]:

```
print("Number of observations (rows):", df1.shape[0])
print("Number of variables (columns):", df1.shape[1])
```

Number of observations (rows): 2626

Number of variables (columns): 21

Python does not give us the initial roles and levels of the variables determined. We have to tag each variable if it is predictor or a target variable. Besides, there is no default setting for python

Q2:

Based on the default settings are variables with missing values rejected? Explain. Based on the default settings, which numeric variables are assigned as nominal variables? Is there a maximum number of levels for character variables in the default settings? Explain.

```
In [74]: # based on the default settings are variables with missing values rejected? Explain  
# this code will display how many missing values are present per column  
print(df1.isnull().sum())
```

```
Att_hrs_fall          0  
GENDER                0  
HIGH SCHOOL PERCENTILE 263  
AGE                   0  
Att_hrs_spr           0  
Avg_income            157  
Distance              98  
Dropped_course        0  
Major_rate             79  
SAT                   0  
Extra_curr             0  
Legacynum              0  
Stu_worker_ind         0  
Need_pct_met           0  
Perc_hrs_comp_fall    0  
Hs_rate                805  
Dorm_rate              0  
Instate                0  
Transcrip              0  
Fall_GPA                0  
Target                  0  
dtype: int64
```

Therefore the code above demonstrates that based on the default settings, variables with missing values are not rejected. The code above proves that most columns have 0 missing values but high_school_percentile, avg_income, distance, major_rate, and Hs_rate all have missing values. Hs_rate has the highest amount, with 805 values missing. Therefore, if a variable has missing values it is not rejected.

In Python, if there are any variables having missing values, the model will drop those missing values itself automatically. However, in SAS, those will be rejected.

In [166...]

```
# is there a maximum number of levels of character variables in the default setting? Explain
# Check data types - to determine which variables are character variables (object)
df1.info()

# this code looks only at the variables that are 'object' and then
# counts the amount of unique levels (categorical variables have a few unique values)
unique_levels = df1.select_dtypes(include='object').agg('nunique')
print(unique_levels)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2626 entries, 0 to 2625
Data columns (total 21 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   Att_hrs_fall    2626 non-null   float64 
 1   GENDER          2626 non-null   object  
 2   HIGH SCHOOL PERCENTILE 2363 non-null   float64 
 3   AGE             2626 non-null   float64 
 4   Att_hrs_spr    2626 non-null   float64 
 5   Avg_income      2469 non-null   float64 
 6   Distance        2528 non-null   float64 
 7   Dropped_course 2626 non-null   object  
 8   Major_rate      2547 non-null   float64 
 9   SAT             2626 non-null   float64 
 10  Extra_curr     2626 non-null   float64 
 11  Legacynum       2626 non-null   object  
 12  Stu_worker_ind 2626 non-null   object  
 13  Need_pct_met   2626 non-null   float64 
 14  Perc_hrs_comp_fall 2626 non-null   float64 
 15  Hs_rate         1821 non-null   float64 
 16  Dorm_rate       2626 non-null   float64 
 17  Instate         2626 non-null   object  
 18  Transcrip      2626 non-null   object  
 19  Fall_GPA        2626 non-null   float64 
 20  Target          2626 non-null   int32  
dtypes: float64(14), int32(1), object(6)
memory usage: 420.7+ KB
GENDER          2
Dropped_course  6
Legacynum       2
Stu_worker_ind  2
Instate         2
Transcrip       2
dtype: int64

```

The code above analyzes the amount of levels in the 7 character variables - shown as 'object'. Most of the character variables have two levels but the Dropped_course variable has 6 levels. Therefore this proves that based on the default setting there is no maximum number of levels for character variables.

The more categories we have in table, the more complex our model will be. Thus, there is no maximum number of levels for character variables. We can create as many categories as we want in python.

Q3:

How many input variables in the DataFrame? How many target variables? How many of the input variables are nominal? Are there any rejected variables? (Compare with your initial answer on the number of variables in this table.) What is a typical value for a sample size? Create a BAR plot of the GENDER variable. What percentage of the students are female? Highlight the students with Target=1 on the above GENDER bar plot. Copy and then paste the plot here. Create a HISTOGRAM plot of the Avg_income variable with 20 bins. Paste here. Take a sample of the last 10 rows of the Retention dataset Take a screenshot (PtrScr on your keyboard) and paste here.

```
In [30]: # How many input variables in the DataFrame?
```

```
input_variables = df1.drop(columns=['Target']) #not include Target variable
num_inp_var = input_variables.shape[1]
print("Number of input variables", num_inp_var)
```

Number of input variables 20

```
In [38]: df1.head()
```

```
Out[38]: Att_hrs_fall  GENDER  HIGH SCHOOL_PERCENTILE  AGE  Att_hrs_spr  Avg_income  Distance  Dropped_course  Major_rate  SAT ...
```

	Att_hrs_fall	GENDER	HIGH SCHOOL_PERCENTILE	AGE	Att_hrs_spr	Avg_income	Distance	Dropped_course	Major_rate	SAT	...	
0	13.0	F		97.0	18.910335	15.0	30573.0	5.083094	0	0.821782	1160.0	...
1	14.0	M		97.0	18.036961	15.0	27305.0	2.198764	0	0.795455	1050.0	...
2	15.0	M		83.0	18.162902	15.0	30573.0	5.083094	0	0.843750	1140.0	...
3	14.0	M		92.0	18.590007	12.0	35865.0	3.245574	0	NaN	1090.0	...
4	16.0	F		80.0	18.880219	16.0	40125.0	19.486585	0	0.840391	1150.0	...

5 rows × 21 columns

```
In [40]: # How many target variables?
```

```
target_variables = df1['Target']

print("Number of target variables", 1)
```

Number of target variables 1

```
In [70]: #Number of nominal variables
nominal_variables = input_variables.select_dtypes(include='object').shape[1] #not include Target variable
print("Number of nominal variables", nominal_variables)
```

Number of nominal variables 6

```
In [98]: # Are there any rejected variables? including missing values and values having 1 unique value
# Columns with all values missing
all_missing = df1.columns[df1.isnull().all()]

# Columns with only one unique value (constant)
constant_cols = [col for col in df1.columns if df1[col].nunique() <= 1]

# Combine
rejected_vars = list(set(all_missing).union(constant_cols))

# Output
print("Number of rejected variables:", len(rejected_vars))
print("Rejected variable names:", rejected_vars)
```

Number of rejected variables: 0

Rejected variable names: []

There are no rejected variables, which is similar to what has mentioned on Q2 before.

```
In [ ]: # What is a typical value for a sample size?
```

Small: < 100 observations — often too small unless it's a very controlled experiment.

Medium: 100–500 observations — usable, but may limit model complexity.

Good: 500–2000 observations — enough for solid statistical analysis and machine learning models.

Large: > 2000 observations — great for complex models and generalization

```
In [106...]: # Create a BAR plot of the GENDER variable. What percentage of the students are female?
import matplotlib.pyplot as plt
import seaborn as sns

# Clean up Gender and Target variables
df1['GENDER'] = df1['GENDER'].astype(str)
df1['Target'] = df1['Target'].astype(int)

# Count how many students in each gender
gender_counts = df1['GENDER'].value_counts()
gender_counts
```


```
Out[106...]: GENDER  
F      1582  
M      1044  
Name: count, dtype: int64
```

```
In [120...]: # Calculate the percentage of students who are female  
total_students = gender_counts.sum()  
female_count = gender_counts.get('F', 0)  
female_percentage = (female_count / total_students) * 100  
print ('the percentage of female students:', female_percentage, '%')
```

```
the percentage of female students: 60.24371667936025 %
```

```
In [130...]: # Create a bar plot of Gender counts  
plt.figure  
gender_counts.plot(kind='bar', color = ['red','lightblue'])  
plt.title('Number of Students by Gender')  
plt.xlabel ('Gender')  
plt.ylabel ('Number of Students')  
plt.xticks(rotation=0)  
plt.show()
```

Number of Students by Gender




```
In [134...]: #Highlight the students with Target=1 on the above GENDER bar plot  
#Count only students with Target = 1 (retained) by Gender  
gender_retained = df1[df1['Target'] == 1]['GENDER'].value_counts()  
gender_retained
```

```
Out[134...]: GENDER  
F    192  
M    142  
Name: count, dtype: int64
```

```
In [144...]: genders = list(gender_counts.index)
```

```
In [158...]: #Plot the stacked bar  
plt.figure  
# base bar  
plt.bar(genders, gender_counts - gender_retained, label = "Target = 0", color = 'lightGray')  
# top bar  
plt.bar(genders, gender_retained, bottom = gender_counts - gender_retained, label = "Target = 1", color = 'skyblue')
```


```
#Chart Labels
plt.title('Student Gender Distribution with Target = 1 highlighted')
plt.xlabel('Gender')
plt.ylabel('Number of Students')
plt.xticks(rotation = 0)
plt.legend(title = 'Retention Status')
plt.show()
```



In [162]: #Create a HISTOGRAM plot of the Avg_income variable with 20 bins

```
plt.figure
plt.hist(df1['Avg_income'].dropna(), bins = 20, color = 'orange', edgecolor='black')
plt.title('Distribution of Average Income')
plt.xlabel('Avg_Income')
plt.ylabel('Number of Students')
plt.tight_layout()
plt.show()
```

Distribution of Average Income



```
In [164]: #Take a sample of the last 10 rows of the Retention dataset  
last_10rows = df1.tail(10)  
last_10rows
```

Out[164...]

	Att_hrs_fall	GENDER	HIGH SCHOOL PERCENTILE	AGE	Att_hrs_spr	Avg_income	Distance	Dropped_course	Major_rate	SA
2616	15.0	M		59.0	18.606434	16.0	36760.0	156.742238	0	0.840391 1290.
2617	12.0	M		13.0	18.880219	18.0	26725.0	167.681727	0	0.795455 1060.
2618	13.0	F		51.0	18.488706	13.0	53190.0	1186.851867	0	0.831585 980.
2619	14.0	M		62.0	18.280630	16.0	60898.0	1154.751207	0	0.831511 1070.
2620	14.0	M		0.6	18.967830	13.0	NaN	NaN	0	0.829268 1140.
2621	13.0	M		42.0	18.151951	12.0	48581.0	179.444678	0	0.838384 1080.
2622	13.0	M		74.0	18.650240	18.0	NaN	110.993592	0	0.831511 1110.
2623	16.0	F		0.6	18.102669	12.0	73510.0	356.674286	0	0.828571 1150.
2624	12.0	F		71.0	18.318960	13.0	NaN	NaN	0	0.829268 1060.
2625	16.0	F		95.0	18.666667	16.0	59599.0	91.022950	0	0.829515 1140.

10 rows × 21 columns

